


Nine attempts to improve the model, measured against the noise bar



No training variant clears the bar. CSLS is the best of a five-value sweep, and multi-view pooling fails a paired McNemar test ($p = 0.13$). The model soup led on validation and then lost on test.